

From Chemistry to History via Kernel Embedding

csaaw

November 16, 2016

abstract

1 introduction

This paper explores the question: how much can the chemical composition of pottery fragments tell us about the evolving connections between settlements in Bronze Age Greece?

Some classical statistical techniques have been used to study the LBNL dataset. For instance, in [main paper], the basic approach consists of using weighted square differences between the chemical composition of two artifacts to determine their similarity, which is then used to cluster the collection of artifacts. This clustering can then be compared to the actual locations of the different sites where the artifacts were found.

The clustering of artifacts according to their chemical composition can say much about the relation between their regions of origin. For instance, if two regions contain many artifacts falling in the same chemical composition cluster, it follows that the materials used in the construction of the artifacts were similar, and perhaps that the two regions shared resources or had markets that were highly connected.

There are some issues with the classical approaches that we wish to attack. Mainly, there is no clear way to analyze connectivity among different archaeological sites based on the clustering of artifacts. A clustering of artifacts based on chemical composition that mimics well the geographical location of the archaeological sites is very illustrative, but it is not obvious how it can be used to quantify the degree of similarity between different sites.

Our contribution to the analysis of the LBNL data is to provide a well-justified similarity measure between archaeological sites based on the chemical composition of their respective artifacts. We are then able to provide quantitative description of the amount of connectivity between different archaeological sites. We used this to classify artifacts into sites and eras based on the chemical features of that artifact. We then checked these predictions about site origin against the actual origin, so as to validate the model.

2 literature review

In this section, two categories of literature are surveyed: (a) theories of Bronze Age trade relevant to the movement of ceramics, and (b) relevant archeometric and modeling methods.

Our dataset includes ceramic shreds from the 7th century BC to the Roman period, but mainly consists of works from the Late Helladic Period - III (LH III, see fig. 1 for a chronology) fig. 2 is a map with dashed lines encoded hypotheses about prominent sea routes in the late Bronze Age, which encompasses LH III.

The evidence for Mycenaean dominance in LHIII primarily comes from Egyptian and Hittite texts, as well as the preponderance of Mycenaean pottery in cities across Greece and the Levant [9].

According to the literature, as more data becomes available and methods for modeling and interpreting the data evolve questions of provenience for specific samples of ware can be traced, differentiations between pottery from the site and imported Mycenaean samples can be made, as well as new insights on relationships between trade sites and the trade routes can be mapped

7000	•	Beginning of widespread use of pottery in Mediterranean. Aegina as maritime center
3200-2000	•	Early Helladic
3200-2500	•	High point of Anatolian trade network
2000-2200	•	Disruption of trade between Cyclades, Mainland, and Crete. Upheaval in Mainland
2000-1550	•	Middle Helladic. Protopalatial and Neopalatial buildings in Crete
1650-1500	•	Late Helladic I / Late Cycladic I / Late Minoan IA
1500-1400	•	Late Helladic II
1400-1300	•	Late Helladic III A
1300-1200	•	Late Helladic III B. High point of Mycenaean influence on trade
1200-1050	•	Late Helladic III C
1050-	•	End of Bronze Age. Transitional Period: Argos, Asine and Berbati rise to prominence.

Figure 1: Timeline of Events of Interest (all dates B.C.)

out. Some work has been made using the LBNL to trace provenience sites as in the work of Grave et al.[3] who argued that focusing on geochemical precincts of regional parent geologies provides a more accurate method for locating provenience. They were able to map out the provenience of Red Lustrous Wheelmade Ware by using two datasets: the NAA results for a sample population of RLW from Bogazkoy and the LBNL Cypriot ceramic NAA dataset. In instances where Mycenaean ceramics are found in other sites in the Mediterranean region such as the work by Jung, Mommsen and Pacciarelli[5] differentiate between vessels originating in Greece versus pottery native to the Punta di Zambrone site.

Maps like fig. 2 are of tremendous historical significance if they can be constructed accurately. A common way to map trade is to identify artifacts that have been transported from their location of origin. How do researchers infer where ceramic objects and shreds originate?


Qualitative methods include categorization by geometric features, and material features. For instance, Davis[1] classifies pieces from LH I Korakou by shapes, as well as surface and burnishing, measured by the eye against a color chart. The task is to sort shreds and pieces by style, and also material - these two elements are necessary to differentiate cases where a style travels but pieces are built from local materials. A qualitative study we'll discuss in detail later is that of Rutter[7] - he uses the shapes and visible properties of pottery found in Korakou to argue that those pieces were erroneously categorized as Mycenaean. That is: qualitative differences are used to prove a classification claim.

Quantitative work in classifying pottery mainly uses data from neutron activation analysis or spectral analysis of paints. This paper uses a neutron activation analysis dataset from the LBNL archeometry archives. The extant work on this dataset (which contains subsections for many geographical areas across the world) has mostly been conducted by Mommsen and colleagues. This work has two main goals: first, to be a proof of concept for quantitative analysis by reproducing known results from limited data, and second, to discover previously unknown connections as a bridge for further inquiry, both qualitative and quantitative. So for instance, in [6], Mommsen et al take their model to succeed by matching existing predictions about early recipe variation in Argolid pottery suggested by Hoffman et al, and by adding new predictions in the case of suggesting a pattern of export from Chania to Cyprus of cream ware. Grave et al[3] follow a similar method of combining the LBNL dataset with other information, also with Cyprus late Helladic wares.

3 Motivation/Dataset

3.1 Motivation and Summary Stats

Archaeological studies has been actively inviting the analytical methods from the entire spectrum of social science studies. However, when it comes to the application of machine-learning techniques, we have not seen much done recently. Therefore, we set out to explore the archaeological databases, searching for datasets that are rich enough, in terms of dimensions of the



Map 7.3 Mycenaean palaces and Aegean maritime routes.

Figure 2: image from [2]




Figure 3: Example of shreds in the database

observables.

Baring the capability of machine-learning in mind, we found the Lawrence Berkeley National Laboratory (LBNL) Nuclear Archaeology Program Archives of particular interest, as it not only contains geographical and historical records of the artifacts, it also carry the chemical compositions for most of its records.

Due to limited time and resources, we focused exclusively on the Greece records, totaling 886 pieces of artifacts from ancient Greece, discovered from 31 archaeological sites¹. The time frame in our sample covers from the Early Helladic era (around 3200 BC) to Roman Republic.

3.2 Data manipulation

For the records of artifacts from Greece, the dataset came with 1,198 observations in the beginning, pulled directly from <http://core.tdar.org/>². The variables include the discovery site, the associating era, geo-coordinates and the chemical compositions of 33 elements³. However, the chemical composition data is not complete. Thus, we dropped a subset of observations as well as insignificant variables of elements (*Sb Ba As Sr V*) to obtain a sample of 886 observations.

4 Problem Formulation

5 problem formulation

Our goal in this project is to understand relations of evolving similarity between archeological sites over time. In this section, we define this problem more exactly.

¹30 in Greece, and one (Perati in Attica) in Turkey.

²To take a quick look at the artifacts, please refer to ??.

³Elements are: *Al Ca V Dy Mn Na K Sr As U Eu Ba Sm La Ti Lu Nd Co Sc Fe Ce Yb Cs Ta Sb Cr Th Ni Rb Tb Hf Zn*

First, a brief intuitive gloss on our technique. In our model, similarity reflects a relationship between our estimates of the underlying functions which are ‘generating’ the artifacts. By modeling it this way, we reduce the site information to an underlying probability function. Then, these functions can be compared, and a relationship computed pairwise between all the sites. This allows us to construct a network. Finally, we can look at the network among sites over the entire period, or segmented by time. In the latter analysis, we separate the artifacts into eras, and look at the similarities in each era separately.

What is this analysis meant to accomplish? Archaeologists have looked at Greek pottery artifacts for centuries in an attempt to understand trade, and each individual site we consider has been described in elaborate detail (for example, see [1]). Our project is meant to be *focused and precise* in its use of data, while aiming at a *general and holistic target* - simplicity of element abundance profile. This might seem like an odd mismatch, but it also has the potential to either turn up novel connections or confirm old results from a fairly independent source.

So the problem we are attempting to solve is to take a well-defined but limited set of features (element abundance) and generate a similarity network using only those features. Then, we will compare the relations in our network to theories and findings from archaeologists and historians about the trade and migration patterns that might effect the similarity of pottery artifacts. In the most general case, we consider the progression of trade power (among our sites, a progression from Aegina in the early Bronze Age, to Festos in the Minoan period, then Mycenae in the late Helladic, and finally Berbati). These theories of trade relations also imply similarity relations - just of a more qualitative kind, and based a far wider body of evidence. Finding significant points of agreement between our model and these theories would be an indication that the similarity relation we’ve calculated based on these simple chemical properties is latching on to something deeper. This could be because of similarities in raw materials, or firing techniques, but most likely a combination of the two. Alternately, finding unexpected similarity relations or relations that run contrary to extant archeological theories might be a call for explanation: could there be a genuine trade connection between these regions? Or if not, could the similarity in composition of the shards be reflection some other kind of connection?

In short, we view the application of machine learning techniques to archeology as aiding in a discovery process. The discovery process in question in this paper is comparing quantitative similarity of chemical composition of pottery to qualitative similarity of regions based on extant theories of influence, migration and trade.

6 Methodology

We are interested in analyzing artifacts from $N = 20$ archaeological sites. We denote by X^i the collection of artifacts in the i -th site, that is $X^i = \{x_1^i, x_2^i, \dots, x_{n_i}^i\}$, where n_i is the number of artifacts in site i and x_m^i refers to the m -th artifact in the i -th site. The chemical composition data available for each artifact consists of the concentration of $L = 27$ different chemicals. Artifact x_m^i then consists of the (normalized) vector containing these chemical concentrations. The simplest way to compare the chemical composition of two artifacts is through their square distance

$$\|x_m^i - x_l^j\|^2 = \sum_{k=1}^L ((x_m^i)_k - (x_l^j)_k)^2$$

where $(x_m^i)_k$ refers to the k -th (normalized) chemical concentration of artifact x_m^i . This in turn is used to construct many similarity measures, for instance

$$k(x_m^i, x_l^j) = \exp(-h\|x_m^i - x_l^j\|^2)$$

where $h > 0$ is a fitting parameter. This corresponds to the widely used Gaussian kernel. Other similarity measures k between artifacts based on their chemical composition are possible. The question of which similarity measure between chemical composition of artifacts is more appropriate for statistical inference is an important one, but we do not consider it further. The challenge is then to construct a similarity measure between different sites based on the similarities

$k(x_m^i, x_l^j)$ between their respective artifacts. We propose constructing such similarity function by looking at the average similarity between all pair of artifacts (x_m^i, x_l^j) belonging to sites X^i and X^j , respectively, given by

$$S(X^i, X^j) = \frac{1}{n_i n_j} \sum_{m=1}^{n_i} \sum_{l=1}^{n_j} k(x_m^i, x_l^j).$$

The justification for this similarity measure between archaeological sites is given by the theory of kernel mean embedding, as described below. We now wish to illustrate how well the similarity measure S performs in finding similarities between archaeological sites based on the chemical composition of their respective artifacts. We do this by providing a network plot of the archaeological sites, where two sites X^i and X^j are connected if their similarity $S(X^i, X^j)$ is larger than a prescribed threshold. The results are illustrated in the figures..

7 Similarity

In this section we detail the way in which similarities are computed. We draw inspiration from a well known set of methods in the machine learning community, and acknowledge that the mathematical machinery might not be familiar to the intended audience. Therefore the reader is invited to skip this section on the first read and come back to it when necessary. We think of the chemical composition vectors $\{x_m^i\}_{m=1}^{n_i}$ corresponding to the artifacts in the i -th site X^i as consisting of samples from some distribution P_i .

7.1 Reproducing Kernels

To build a network we need an appropriate notation of similarity between sites (probability densities). We desire the probability densities to inhabit a well behaved space of functions (a Hilbert Space) where a common notion of similarity exists. We propose to embed each pd into the so called RKHS of a kernel K and use inner product in this space as the similarity measure. The similarities we propose capture similarities among probability densities as elements of a function space. To determine these similarities we need first determine the function space the probability densities inhabit. The function space we will use is the so called Reproducing Kernel Hilbert Space (RKHS) associated to a reproducing kernel. We now define these mathematical objects. A *reproducing kernel* is a symmetric function $k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$ for which there exists a unique Hilbert space of functions \mathcal{H} such that $k(\cdot, x) \in \mathcal{H}$ for all $x \in \mathcal{X}$ and the reproducing property

$$f(x) = \langle f, k(\cdot, x) \rangle \quad (1)$$

holds for all $f \in \mathcal{H}$ and all $x \in \mathcal{X}$

7.2 Kernel Mean Embedding

Let k be a reproducing kernel as in section, the *kernel mean embedding* of the probability density P in the RKHS \mathcal{H} of k is

$$\phi_0(P) = \int k(\cdot, x) P(x) dx$$

Since the true form of P is unknown we use the available data to estimate the KME. Let $x_{l=1}^n$ be an iid sample from P , then the empirical kme of O is

$$\phi(P) = \frac{1}{n} \sum_{l=1}^n k(\cdot, x_l).$$

Although $\phi(P)$ depends on the sample we will omit this dependencies to simplify notation. With this tool at hand, we define the similarity between site i and site j as the inner product of their

corresponding KME's.

$$\begin{aligned}
S(X^i, X^j) &= \langle \phi(P_i), \phi(P_j) \rangle_H \\
&= \frac{1}{n_i n_j} \sum_{l=1}^{n_j} \sum_{m=1}^{n_j} \langle k(\cdot, X_l^i), k(\cdot, x_m^j) \rangle_H \\
&= \frac{1}{n_i n_j} \sum_{l,m} k(x_l^i, x_m^j)
\end{aligned}$$

where the last equality follows from the reproducing property.

8 Results/explanations

In fig. 4, we generated similarity plots for the various sites over five eras in order to see how well our groupings based on chemical composition map on to connections posited by historians and archeologists based on a much larger set of data. Because our data comes from a fairly long historical period, we needed to break it up into eras in order to have a well-defined set of historical theories and connections with which to compare it. That is, historians are not generally in the business of theorizing about the connection between Mycenae and Palaikastro over the entire bronze age, but over much smaller periods.

We chose periods 2-4 based on the standard chronology for bronze age Greece displaying in fig. 1: period 2 (1399-1300) is the Late Helladic III-A, period 3 (1299-1200) is the Late Helladic III-C, and period 4 (1199-1050) is the Late Helladic III-C. The first period includes the earliest artifacts in the dataset; we allowed a much larger period because there were fewer sites and artifacts from this era, and likewise with the final period, which includes artifacts up to Roman times.

How many sites in each period?

In fig. 4, every site with artifacts dating from that period was included in the graph for that period. The blue links show a kernel-based similarity in the sites based on those artifacts, as described in SECTION. Every similarity measure over a threshold is represented. We chose this threshold based on the histogram of the sum number of connections according to each possible threshold: 10^5 was selected since it was near the median across all eras.

One way of analyzing these plots is to look at the most connected site in each era, that is, the site with the most number of links. A link between sites could represent either geographic factors or cultural-historical ones. Geographic factors include similarity in raw materials available at those sites, either due to proximity or similarity in geographic features such as altitude or distance from bodies of water. Cultural-historical factors include trade, influence and other forms of resettling, conflict and movement of people. For instance, artifacts could reflect trade relations in their chemical composition by actually sharing an origin if they were produced in one site and traded to another, or more indirectly by reflecting similar styles in material, such as a trend in red ware that spread via cultural influence rather than the literal trade in red ware pots. Note that many stylistic features would not be represented in chemical composition, such as drawings and figures which are often used to track influence [4]. Thus, the most connected site is an indication of some kind of centrality or similarity, but a priori it is hard to predict what kind of similarity would show up in chemical composition.

The first thing to notice about fig. 4 is that the connection between sites change significantly throughout the eras. In fact no two recurring sites retain a link in all graphs where they both appear. Of course, geographical proximity between sites does not change over time, nor does altitude, or proximity to water and natural resources (for the most part). This alone indicates that our similarity measure is not determined by these geographic factors.

Starting in era 1, we see that Aegine is the most connected site, linked to 5 others whereas the next most connected sites are only linked to 2 others. Aegine was a powerful city in the middle Bronze age, so this is not surprising - though there are relatively few artifacts and a long period here so not much can be safely inferred.

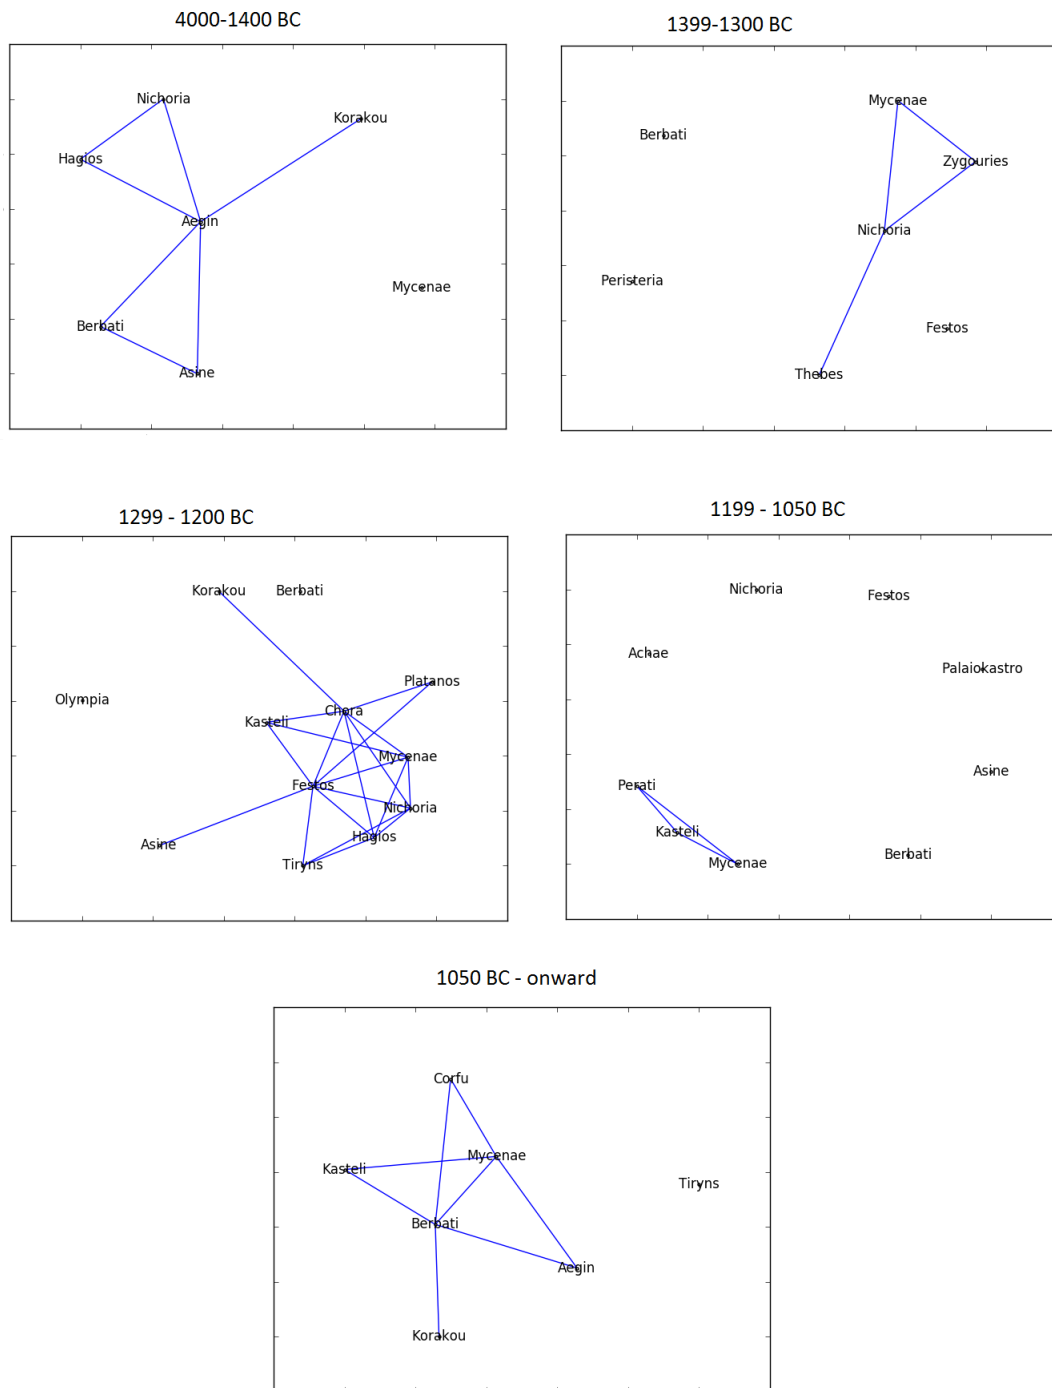


Figure 4: Evolution of networks of the Archaeological sites

Network using all artifacts




Figure 5: Overall network of the Archaeological sites

In era 2, the Late Helladic III-A, Nichoria is the most connected, with 3 links, followed by Zygouries and Mycenae with 2. Nichoria is already theorized to have been particularly prominent in III-A and III-B[8], and even though it has artifacts represented in era 4, it does not have any links. The entire late Helladic-III is thought to be characterized by Mycenaean dominance[2], and it is in this period that Mycenae starts to be among the most connected nodes in our graph.

Era 3 encompasses the Late Helladic III-B period, the high period of the Mycenaean trade network [2]. As expected, there is both higher overall connectivity, and Mycenae is among the most connected nodes.

Era 4 has less connectivity, even relative to the number of sites represented, with only a triad of Perati-Mycenae-Kasteli Chania. This could be indirectly due to the fact that this marks the beginning of the dark age. A reduction in economic power and trade could lead to fewer relations between cities and thereby more local cohesion. Alternately, reduction in trade could reduce cultural influence and thereby the similarity of materials used in pottery might have diverged.

Era 5 again shows a cluster containing Mycenae towards the center, along with Berbati. It's hard to draw historical conclusions over such a long timespan, but Berbati

Overall,

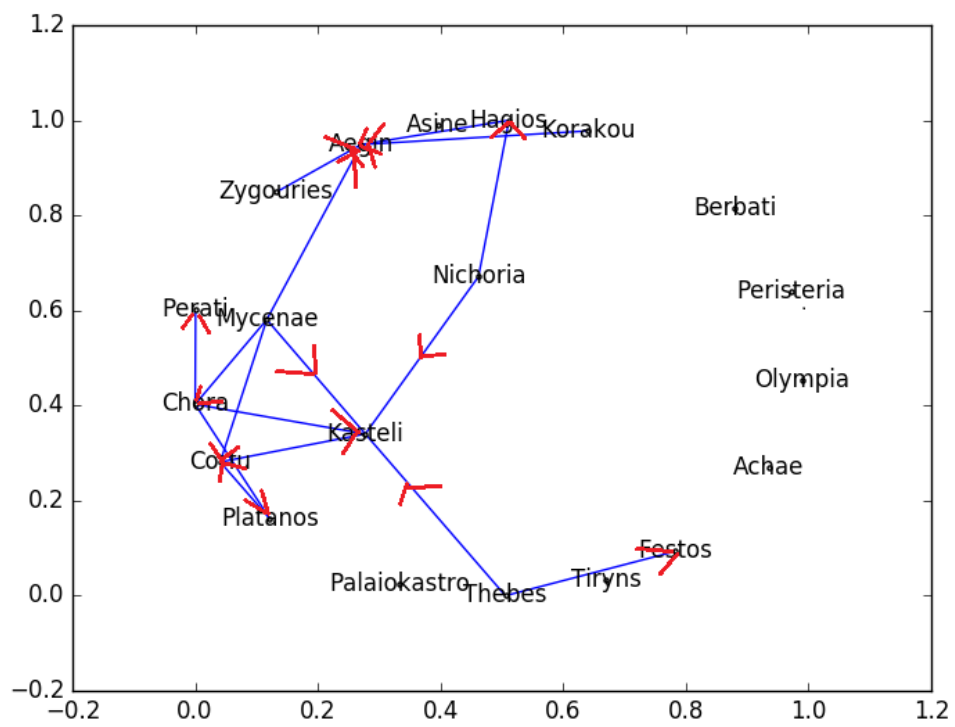


Figure 6: Directed network of the Archaeological sites

9 Conclusion

10 method

References

- [1] Jack L Davis. “Late Helladic I Pottery from Korakou”. In: *Hesperia: The Journal of the American School of Classical Studies at Athens* 48.3 (1979), pp. 234–263.
- [2] Nancy H Demand. *The Mediterranean context of early Greek history*. John Wiley and Sons, 2011.
- [3] Peter Grave et al. “Ceramics, trade, provenience and geology: Cyprus in the Late Bronze Age”. In: *Antiquity* 88.342 (2014), pp. 1180–1200.
- [4] Julie Hruby. “Mycenaean pottery from Pylos: an indigenous typology”. In: *American Journal of Archaeology* (2010), pp. 195–216.
- [5] Reinhard Jung, Hans Mommsen, and Marco Pacciarelli. “From west to west: Determining production regions of Mycenaean pottery of Punta di Zambrone (Calabria, Italy)”. In: *Journal of Archaeological Science: Reports* 3 (2015), pp. 455–463.
- [6] H Mommsen, Th Beier, and A Hein. “A complete chemical grouping of the Berkeley neutron activation analysis data on Mycenaean pottery”. In: *Journal of Archaeological Science* 29.6 (2002), pp. 613–637.
- [7] Jeremy B Rutter. “Ceramic evidence for northern intruders in Southern Greece at the Beginning of the Late Helladic IIIC Period”. In: *American Journal of Archaeology* (1975), pp. 17–32.
- [8] Cynthia W Shelmerdine. “Nichoria in context: a major town in the Pylos kingdom”. In: *American Journal of Archaeology* 85.3 (1981), pp. 319–325.
- [9] Malcolm H Wiener. “The absolute chronology of Late Helladic III A2 revisited”. In: *The Annual of the British School at Athens* 98 (2003), pp. 239–250.